




Reinforcement Learning

Slides by Ekapol Chuangsuwanich and
Nat Dilokthanakul, researcher at VISTEC


3 Modes of Learning




Supervised Learning



Reinforcement Learning




Unsupervised Learning



3 Modes of Learning

Supervised Learning




- Observe:
 - $(x_1, y_1), (x_2, y_2), \dots$
- Objective:
 - Input an unseen x_{new}
 - What is y_{new} ?



3 Modes of Learning

Unsupervised Learning

- Observe:
 - $x_1, x_2, x_3, x_4, \dots$
- Objective:
 - What is $P(x)$?
 - What is a *good* representation of x ?
 - What can we learn from $P(x)$?

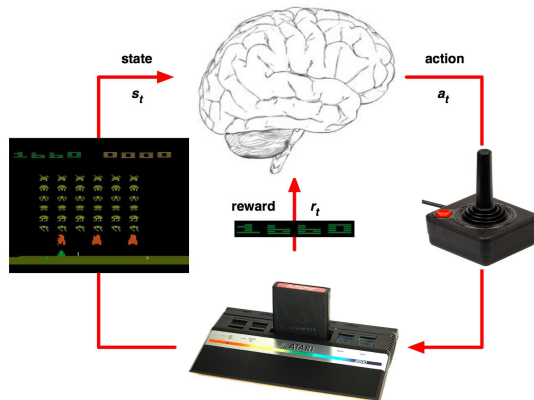


3 Modes of Learning

Reinforcement Learning (RL)



- Observe:
 - The states (x_1, x_2, x_3, \dots)
 - The reward (r_1, r_2, r_3, \dots)
- Can also take actions
 - a_1, a_2, a_3, \dots
- What are the best actions?
 - Such that we will receive highest accumulative rewards



Applications

- Robotic
- Games
- Cooling system
- Autonomous vehicle
- etc.

DeepMind AI Reduces Google Data Centre Cooling Bill by 40%

<https://deepmind.com/blog/deepmind-ai-reduces-google-data-centre-cooling-bill>
-40/



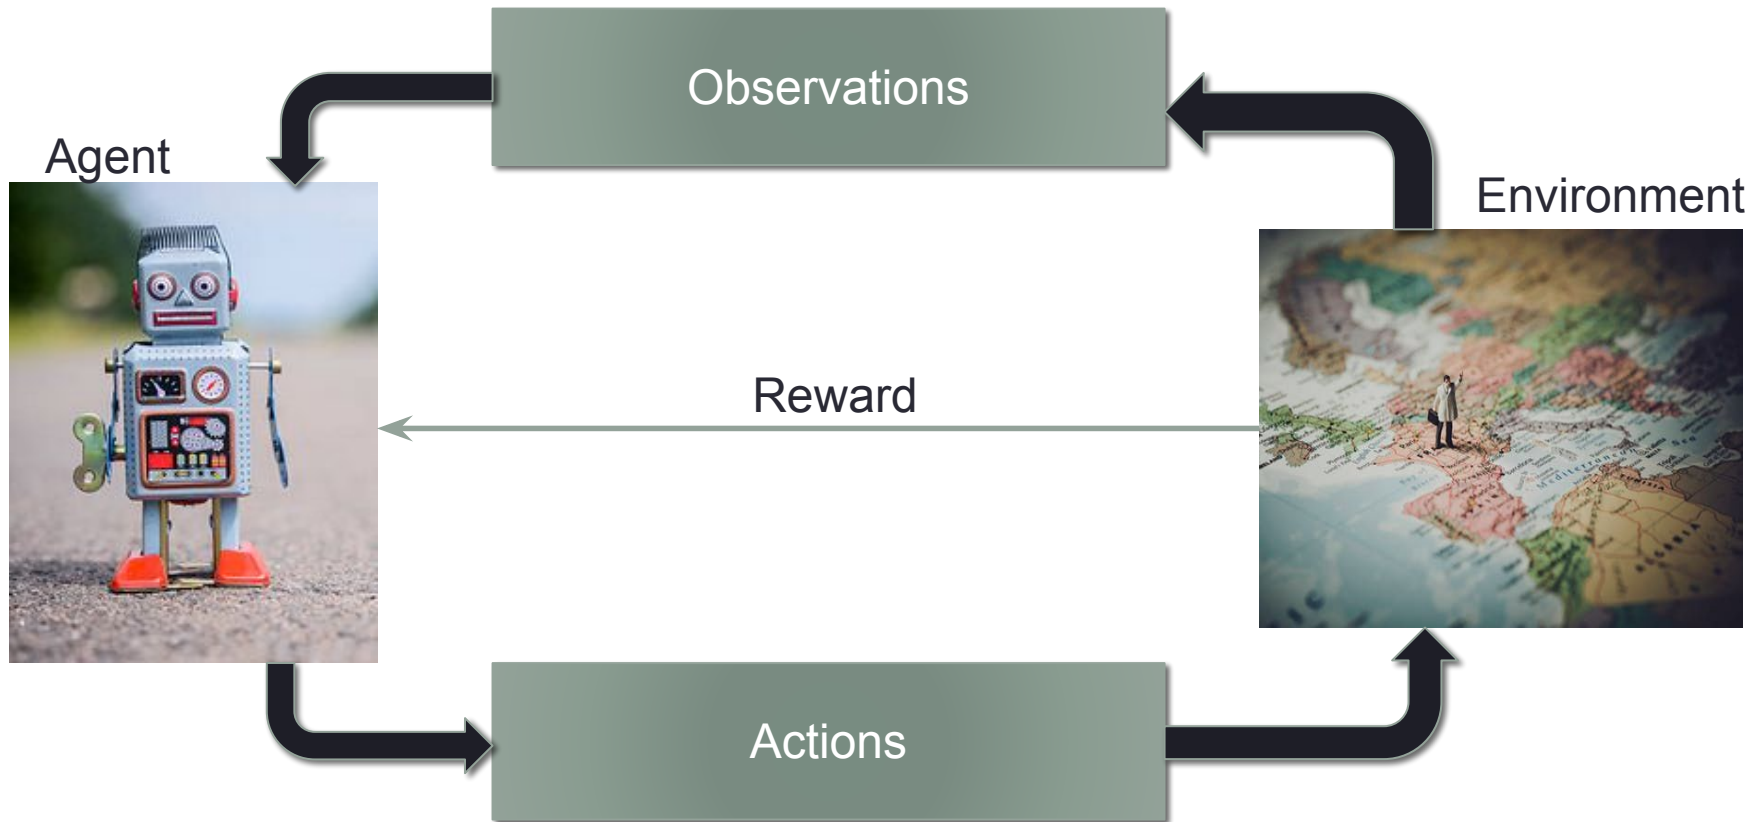
<https://research.googleblog.com/2016/03/deep-learning-for-robots-learning-fro>



<http://spectrum.ieee.org/automaton/robotics/drones/drone-uses-ai-and-11500-crashes-to-learn-how-to-fly>



RL framework




Learning through trial and error

RL framework

Reward (r_t)

State (s_t)

Action (a_t)



Markov property

$$a_t = A(s_t)$$

$$r_{t+1} = R(s_t, a_t)$$


$$s_{t+1} = S(s_t, a_t)$$

Rewards-based learning

- Maximise the rewards
- Can we design any desired behaviour with reward?



$$R_t = \Delta \text{distance}$$



$$R_t = \text{score}$$



$$R_T = \begin{cases} 1, & \text{win} \\ -1, & \text{lose} \end{cases}$$

The Environment


How can we model the environment?

Markov Decision Process (MDP)

- **S, A, P, R, γ**
- **S** – Set of states
- **A** – Set of actions
- **P** – Transition between states given an action

$$P_{s,s'}^a = \text{Prob}[s_{t+1}=s' \mid s_t=s, a_t=a]$$

- **R** – Rewards associated with actions and states
- **γ** – Discount factor




Markov Decision Process (MDP)

Environment could be stochastic
One action, multiple outcomes

- **S, A, P, R, γ**
- **S** – Set of states
- **A** – Set of actions
- **P** – Transition between states given an action

$$P_{s,s'}^a = \text{Prob}[s_{t+1}=s' \mid s_t=s, a_t=a]$$

- **R** – Rewards associated with actions and states
- **γ** – Discount factor



Markov Property

$$p(s_{t+1} | s_t, a_t)$$


- s_{t+1} depends only on s_t
- not s_{t-1} , not anything before
- this simplifies our situation!

But, is it true in every case?

- It depends on your observed state
 - Fully observable state
 - Partially observable state



Fog of war




Fully observable

Fully Observable State

- Fully observable state: All information from the past is captured in the current state



For Go, a board position
For simple video games, stack
multiple frames



The Agent

Policy

- Policy = a mapping from a state to an action
- Objective of RL is to find the “*optimal*” policy!


$$a = \pi(s)$$

Can be either
deterministic or
stochastic

$$a \sim \pi(s)$$

Policy

Example: tabular policy




Learning

How do we find the best policy ?

Rollout (Our data)

Time	0	1	2	3	...	T-1	T
S	S_0	S_1	S_2	S_3	...	S_{T-1}	S_T
A	A_0	A_1	A_2	A_3	...	A_{T-1}	A_T
R	R_0	R_1	R_2	R_3	...	R_{T-1}	R_T
Done	0	0	0	0	...	0	1

Don't care



Return (Cumulative rewards)

- Return = **cumulative** rewards with discount

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots = \sum_{t'=t}^{\infty} \gamma^{t'-t} r_{t'}$$



$r_t = 0$



$r_{t+10} = 1$



$r_{t+34} = -1$

What is learning?

- Use data to find/search for the best policy


What is the best policy?

- Policy that give us the highest expected return!

$$G = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$$

$$J(\pi) = E_{\pi}[G]$$

Return under a policy (G^π)



* THIS IS JUST ONE TRY!
ANOTHER MIGHT YIELD DIFFERENTLY

Expected Return

How good “on average” is our return?

Statistical expectation

For value G_0

Following policy π

$E_{\pi} [G_0]$

- Start
- Play with π
- $G_0 = \sum_{t=0}^{\tau} r_t$
- Retry many times (∞)
- Average G_0

A naive learning method

1. Initialise a policy randomly
2. Evaluate the policy by running that policy multiple times
 - a. which we then collect the returns of all the runs
3. Randomly initialise another policy
4. Evaluate the new policy
5. Keep the policy that have a higher expected return
6. Repeat 3-5

Intuitive. But very inefficient!

How to make it more efficient?

Model-free RL

Q-learning algorithm

- Let's define a state value as
 - Expectation of the return after visit s and follow π

$$V^{\pi}(s) = E_{\pi}[G_t | s_t = s]$$

- Let's define a state-action value (Q-value) as
 - Expectation of the return after visit a state s , take action a

$$Q^{\pi}(s, a) = E_{\pi}[G_t | s_t = s, a_t = a]$$

$$V^{\pi}(s) = E_{\pi(s)}[Q^{\pi}(s, \pi(s))]$$

Q-learning algorithm

- There exist an optimal value function associate with an optimal policy,

$$V^*(s) = \max_{\pi} V^{\pi}(s) \quad \forall s \in \mathcal{S}$$

- The optimal policy is the policy that achieves the highest value for every state

Q-learning algorithm

- It follows that

$$V^*(s) = \max_a [Q^*(s, a)]$$

- and ..

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$


- Optimal actions can be found indirectly through Q-value

Example, Tabular Q-learning

			+1
			-1

Monte-Carlo Estimator


- Initialise π



Monte-Carlo Estimator


$$Q^\pi(s, a) \approx r_0 + \gamma r_1 + \gamma^2 r_2 + \dots + \gamma^n r_n$$

$$\gamma = 0.9$$




Policy Improvement

$$\pi'(s) = \operatorname{argmax}_a Q(s, a)$$




Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$




Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$




Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$




Bootstrap Estimator

$$Q(s, a) \approx r_0 + \gamma \max_b Q(s', b)$$



Policy Improvement



Bias and Variance in RL

What is bias of V_π estimation?

- Let $\hat{V}_\pi(s)$ be an estimate of $V_\pi(s)$
- $\hat{V}_\pi(s)$ is unbiased if:

$$\mathbb{E}_s \left[\hat{V}_\pi(s) - V_\pi(s) \right] = 0$$

What is variance of $\hat{V}_\pi(s)$ estimation?

- $\text{Var} \left[\hat{V}_\pi(s) \right] = \mathbb{E}_s \left[(\hat{V}_\pi(s) - \mathbb{E}_s \left[\hat{V}_\pi(s) \right])^2 \right]$
- High if $\hat{V}_\pi(s)$ fluctuates a lot

Bias and Variance

- Monte-Carlo estimate has high variance and low bias.
- Bootstrap estimate has higher bias but lower variance.

Function Approximator (FA)


- Tabular Q-value is impractical when the state-action space is large!
 - Need large memory
 - Impractical to fill up every cell
- Enter .. a function approximator

Function Approximator (FA)

- Tabular Q-value is impractical when the state-action space is large!
 - Need large memory
 - Impractical to fill up every cell
- Enter .. a function approximator

Function Approximator (FA)

- Instead of a table containing Q-value for every state and action, use a function that output Q-values.



Learning with FA

- With tabular Q-learning,
 - the act of learning = putting Q-value in the table
- With function approximator,
 - the act of learning = searching for the optimal parameters of the FA

Learning with FA

- How to adapt the parameters (weights) of the FA?
- Step 1: Define a loss function.
- Step 2: Optimise the weights to minimise the loss

Loss function

- What should be the loss function?
- Introducing Bellman's equations

$$V^{\pi}(s_t) = E_{\pi, P}[r_t + \gamma V^{\pi}(s_{t+1})]$$

$$Q^{\pi}(s_t) = E_{\pi, P}[r_t + \gamma Q^{\pi}(s_{t+1}, a_{t+1})]$$

- Bellman's optimality equations

$$Q^*(s_t) = E_P[r_t + \gamma \max_b Q^*(s_{t+1}, b)]$$

Loss function

- The Bellman's equation must hold for correct Q-value
- Rewrite the Bellman's optimality with our estimator (FA)
-

$$\hat{Q}(s_t) = E_P[r_t + \gamma \max_b \hat{Q}(s_{t+1}, b)]$$


**The estimator is correct if
the left hand side = right hand side**

$$TD = r_t + \gamma \max_b \hat{Q}_\theta(s_{t+1}, b) - \hat{Q}_\theta(s_t)$$

“Temporal Difference error”

Temporal Difference Learning

- Use TD-error to guide learning
- Example
 - Deep Q-Networks (DQN)
 - Deep convolutional neural network as a function approximator
 - Optimise square TD-error



$$L(\theta) = (r_t + \gamma \max_b \hat{Q}_\theta(s_{t+1}, b) - \hat{Q}_\theta(s_t))^2$$

Policy Gradient methods

Policy gradient

Q Learning

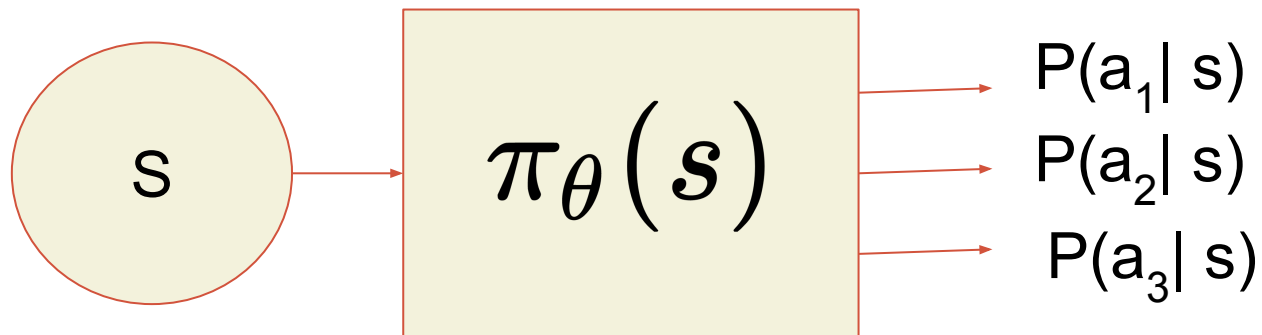

- policy is implicit
- if we already have Q, we have policy
- we just look at Q to get π

Policy gradient

- learns π directly explicitly
- use Q, V as a helper for learning π

Policy gradient

- Use Function Approximator to represent policy directly



Loss function

- Start-state objective

$$J(\theta) = E_{\pi(\theta)} [G | s_0] = V(s_0)$$

- Average-reward objective

$$J(\theta) = \sum_s d^\pi(s) \sum_a \pi(s, a) r(s, a)$$

* \mathbf{d} is a stationary distribution of a Markov chain.

- One way to optimise these objectives is to use SGD.

Computing the gradient

Let's try to compute the gradient of the start-state objective

$$J(\theta) = E_{\pi_{\theta}} [G | s_0]$$

To evaluate this expectation, maybe we could try a one-sample Monte-Carlo estimator:

$$J(\theta) \approx r_0 + \gamma r_1 + \gamma^2 r_3 + \dots$$

$$\nabla J(\theta) \approx \nabla_{\theta} [r_0 + \gamma r_1 + \gamma^2 r_3 + \dots] \quad \times$$

Doesn't quite work? The evaluated value does not depend on θ . Gradient can't be computed.

Computing the gradient

Maybe we can try change θ a little bit and find the difference?

$$J(\theta) \approx r_0 + \gamma r_1 + \gamma^2 r_3 + \dots$$

$$J(\theta + \delta\theta) \approx r'_0 + \gamma r'_1 + \gamma^2 r'_3 + \dots$$

$$\nabla J = \frac{J(\theta) - J(\theta + \delta)}{\delta}$$

Could work? But...

Looks very expensive and noisy to compute!

Maybe there is a better way?

Policy gradient

Let's start from the average-reward objective

$$J(\theta) = \sum_s d^\pi(s) \sum_a \pi_\theta(s, a) r(s, a)$$

For simplicity let's assume $d(s)$ does not depend on θ

$$J(\theta) = \sum_s d(s) \sum_a \pi_\theta(s, a) r(s, a)$$

$$\nabla_\theta J(\theta) = \sum_s d(s) \sum_a \nabla_\theta \pi_\theta(s, a) r(s, a)$$

Almost there...

Policy gradient

REINFORCE trick!

$$\nabla_{\theta} J(\theta) = \sum_s d(s) \sum_a \nabla_{\theta} \pi_{\theta}(s, a) r(s, a)$$

$$\nabla_{\theta} J(\theta) = \sum_s d(s) \sum_a \pi_{\theta} \nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)$$

$$\nabla_{\theta} J(\theta) = E_{\pi} [\nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)]$$

$$\nabla_{\theta} J(\theta) \approx \nabla_{\theta} \log \pi_{\theta}(s, a) r(s, a)$$

Policy gradient theorem

There is a theorem...called policy gradient theorem
say that we can replace $r(s, a)$ with $Q(s, a)$


$$\nabla_{\theta} J(\theta) = E_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^{\pi_{\theta}}(s, a)]$$

Encourage Exploration


- policy $\pi_\theta(a|s)$ could be too confident early
- Like, $\pi_\theta(a = a|s) = 1$
- this could lead to insufficient exploration
- encourage exploration by “entropy term” $H(\pi_\theta)$
- we want to punish too low entropy
- New gradient rule: $\nabla_\theta J(\theta) \rightarrow \nabla_\theta J(\theta) + \nabla_\theta H(\pi_\theta)$

Entropy (H)

high entropy



low entropy



On policy and off policy algorithms

Q Learning: $Q^*(s_t, a_t) = r_{t+1} + \max_a Q^*(s_{t+1}, a)$

- you need $a_t, s_t, r_{t+1}, s_{t+1}$ to satisfy the above equation
- you can get (a, s, r, s') from any policy
- Q learning is *off-policy*

Policy gradient: $\nabla_{\theta} J(\theta) = \mathbb{E} [Q_{\pi}(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)]$

- you need s, a and $Q_{\pi}(s, a)$
- Q_{π} needs to be from the current policy
- s, a needs to come from current policy
- policy gradient is *on-policy*


Baselines

$Q_\pi(s, a)$ has high variance (Monte Carlo)


$Q_\pi(s, a) \nabla_\theta \log \pi_\theta(a|s) = \nabla_\theta J(\theta)$ is very noisy

Slow down training a lot!

We need to reduce variance to speed up the training



Baselines reduce variance



What is a good $b(s)$?

$V_\pi(s)$ is a convenient choice

Why baseline use $b(s)$ not $b(s, a)$

- $b(s, a)$ has potential to reduce more variance, but harder to incorporate to the policy gradient theorem
- $b(s)$ can be added without changing the objective

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{s \sim d^{\pi}, a \sim \pi} [(Q_{\pi}(s, a) - b(s)) \nabla_{\theta} \log \pi_{\theta}(a|s)] \\ &= \mathbb{E} [Q_{\pi}(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)] - \mathbb{E}_{s \sim d^{\pi}, a \sim \pi} [b(s) \nabla_{\theta} \log \pi_{\theta}(a|s)]\end{aligned}$$

Consider:

$$\begin{aligned}&\mathbb{E}_{s \sim d^{\pi}, a \sim \pi} [b(s) \nabla_{\theta} \log \pi_{\theta}(a|s)] \\ &= \mathbb{E}_{s \sim d^{\pi}} \left[\int_a \pi(a|s) b(s) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi(a|s)} \right] \\ &= \mathbb{E}_{s \sim d^{\pi}} [b(s) \nabla_{\theta} \int_a \pi_{\theta}(a|s)] \\ &= \mathbb{E}_{s \sim d^{\pi}} [b(s) \nabla_{\theta} 1] = 0\end{aligned}$$

$b(s)$ doesn't affect the objective

Advantage Function

When we use $V(s)$ as baseline:

$$A(s, a) = Q(s, a) - V(s)$$

We call this the **advantage function**.

It tells the relative value of the actions.

Lower variance than using absolute value of the actions.

Model-based RL

Model-based RL

- In model-based RL, we first build the model of the environment
- Then use that model to directly search for the answer.
- The problem is ... inaccurate model can give us bad policies...
- It is believed that if we can treat the uncertainty in the model correctly...model-based RL is the most efficient method!
- However, measuring uncertainty in the model is also very difficult.

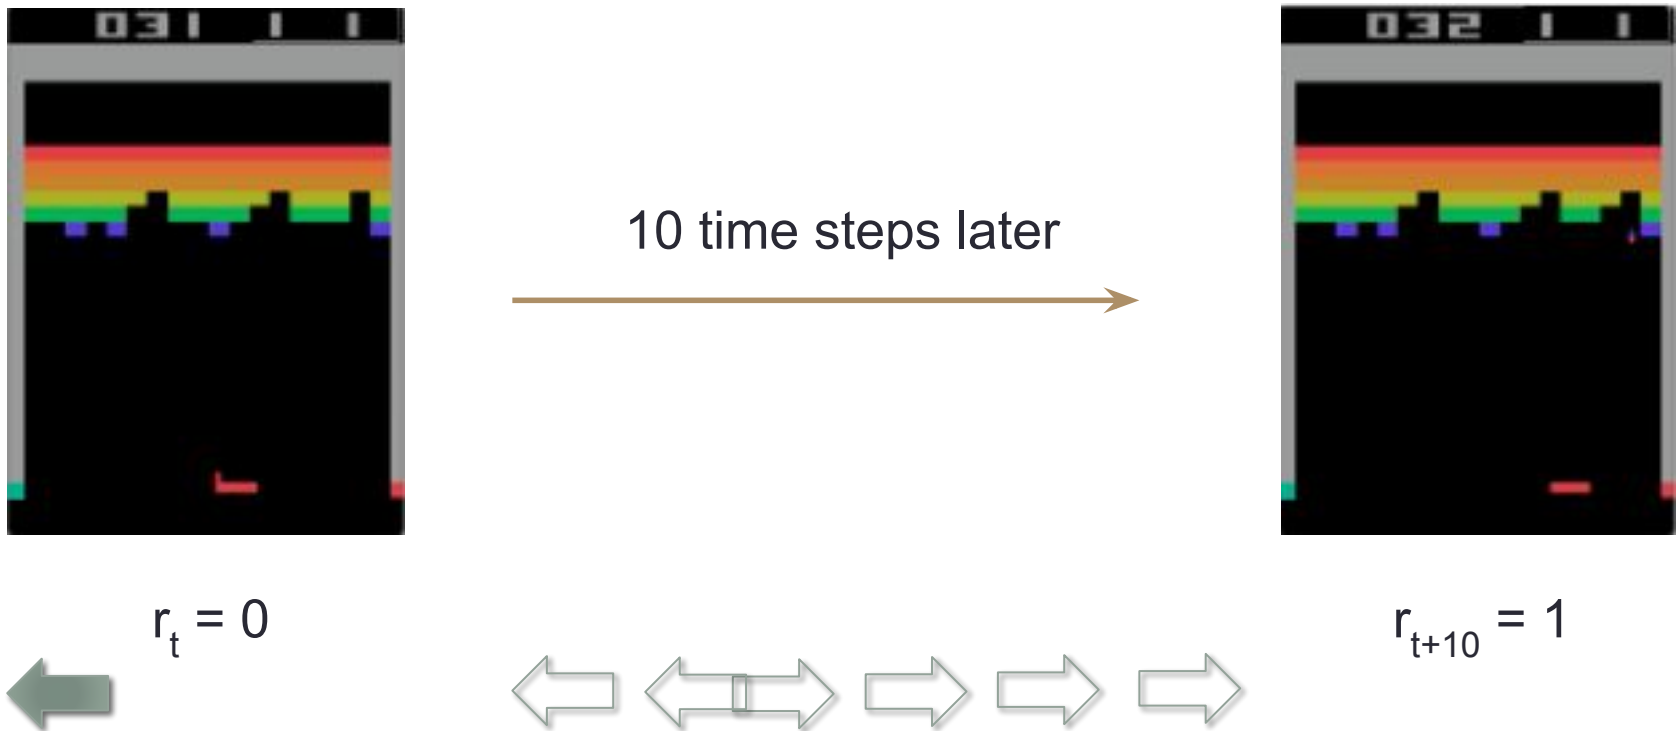
Things to consider

When do we need RL?

- Your action affects the observation.
 - $x_1, a_1 \longrightarrow x_2$
- The target behaviour is difficult to be directly hard-coded.
- Collection of the data of a target behaviour is difficult.

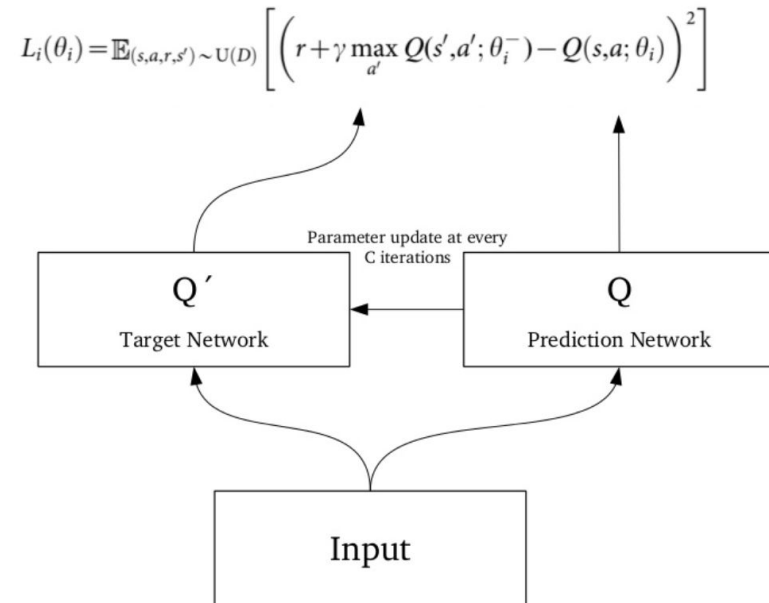
Credit assignment problem

- An action can have consequences further away in time
- Some movements might not have any effect on the outcome



Bias and Variance

- Bias and variance are really important in RL
- We want to reduce both of them as much as possible
 - DQN uses experience replay to reduce bias
 - DQN uses target network to reduce variance
 - Actor-Critic method uses baseline to reduce variance
 - Actor-Critic method uses parallel worker to reduce bias
 - etc.



<https://www.slideshare.net/MuhammedKocaba/human-level-control-through-deep-reinforcement-learning-presentation>

Exploration and Exploitation

- Many RL results assume that all of the states are visited infinitely often.
- Also, many RL algorithms are reduced into just an optimisation problem.
- Therefore, nicely spread/informative data can help a lot!
- DQN uses epsilon-greedy for exploration
- Policy gradient uses entropy regulariser to encourage exploration

Sparse reward problem

- Another problem is when rewards are sparse.
- Since, model-free RL is just learning the correlations of trajectories and rewards... when there is no reward, RL cannot learn.
- Can we make it better?
 - Curiosity + intrinsic motivation?
 - Curriculum learning?
 - Hierarchical RL?

Optimisation problem

- Initialization problems
- Is SGD the best we can do?
- Natural gradient?
- Adaptive learning rate?
- Catastrophic forgetting?

Designing the reward signal

- Reward design can be quite challenging..
- Naive reward design can lead to unexpected (cheating) behaviours!
- Example:



Current trends & open problems

- Intrinsic motivation, reward-bonus
- Imitation learning
- Multi-agent system and self-play
- Curriculum learning
- Model-based RL
- Robot learning + sim-to-real transfer learning
- etc...

Reinforcement learning

Elements of RL

- Environment, Agent, State, and MDP

Estimating Q

- Monte-Carlo, Bootstrap

- Deep learning as a function approximator

Policy learning

- Q-learning

- TD learning

- Policy gradient

Concepts

- Exploration vs Exploitation